

sales data analysis

```
In [26]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score
```

```
In [3]: df = pd.read_csv("sales_data.csv")
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  200 non-null   object
1   Invoice_ID             200 non-null   object
2   Product_Name          200 non-null   object
3   Category              200 non-null   object
4   Quantity              200 non-null   int64
5   Price                 200 non-null   int64
6   Total_Sales           200 non-null   int64
7   Customer_ID           200 non-null   object
8   Payment_Method        200 non-null   object
dtypes: int64(3), object(6)
memory usage: 14.2+ KB
```

```
In [5]: df.head()
```

Out[5]:

	Date	Invoice_ID	Product_Name	Category	Quantity	Price	Total_Sale
0	2024-01-01 00:00:00.000000000	INV1000	Smartphone	Electronics	1	49098	49098
1	2024-01-01 21:49:44.924623115	INV1001	Tablet	Electronics	2	48765	97530
2	2024-01-02 19:39:29.849246231	INV1002	Jeans	Fashion	4	460	1840
3	2024-01-03 17:29:14.773869346	INV1003	Tablet	Electronics	2	33618	67236
4	2024-01-04 15:18:59.698492462	INV1004	T-shirt	Fashion	2	4664	9328

```
In [6]: df.describe()
```

Out[6]:

	Quantity	Price	Total_Sales
count	200.000000	200.000000	200.000000
mean	3.015000	12102.030000	36222.185000
std	1.379962	14758.015052	50919.852045
min	1.000000	229.000000	229.000000
25%	2.000000	2264.250000	4767.500000
50%	3.000000	3878.500000	14280.000000
75%	4.000000	19248.250000	43772.500000
max	5.000000	49997.000000	249985.000000

In [7]:

df.columns

Out[7]:

Index(['Date', 'Invoice_ID', 'Product_Name', 'Category', 'Quantity', 'Price', 'Total_Sales', 'Customer_ID', 'Payment_Method'], dtype='object')

In [8]:

df.isnull().sum()

Out[8]:

Date0
Invoice_ID0
Product_Name0
Category0
Quantity0
Price0
Total_Sales0
Customer_ID0
Payment_Method0
dtype: int64

In [9]:

df.drop_duplicates(inplace=True)

In [10]:

df.head()

Out[10]:

	Date	Invoice_ID	Product_Name	Category	Quantity	Price	Total_Sales
0	2024-01-01 00:00:00.000000000	INV1000	Smartphone	Electronics	1	49098	49098
1	2024-01-01 21:49:44.924623115	INV1001	Tablet	Electronics	2	48765	97530
2	2024-01-02 19:39:29.849246231	INV1002	Jeans	Fashion	4	460	1840
3	2024-01-03 17:29:14.773869346	INV1003	Tablet	Electronics	2	33618	67236
4	2024-01-04 15:18:59.698492462	INV1004	T-shirt	Fashion	2	4664	9328

Exploratory Data Analysis (EDA)

```
In [11]: total_revenue = df["Total_Sales"].sum()
        print("Total Revenue:", df["Total_Sales"].sum())

Total Revenue: 7244437

In [12]: total_invoices = df["Invoice_ID"].nunique()

In [13]: total_customers = df["Customer_ID"].nunique()

In [14]: best_products = df.groupby("Product_Name")["Total_Sales"].sum().sort_values(ascending=False)

In [15]: category_sales = df.groupby("Category")["Total_Sales"].sum().sort_values(ascending=False)

In [16]: payment_counts = df["Payment_Method"].value_counts()

In [17]: df["Month"] = pd.to_datetime(df["Date"]).dt.month_name()
        monthly_sales = df.groupby("Month")["Total_Sales"].sum()

In [18]: top_customers = df.groupby("Customer_ID")["Total_Sales"].sum().sort_values(ascending=False)
        df.head(5)
```

Out[18]:

	Date	Invoice_ID	Product_Name	Category	Quantity	Price	Total_Sales
0	2024-01-01 00:00:00.000000000	INV1000	Smartphone	Electronics	1	49098	49098
1	2024-01-01 21:49:44.924623115	INV1001	Tablet	Electronics	2	48765	97530
2	2024-01-02 19:39:29.849246231	INV1002	Jeans	Fashion	4	460	1840
3	2024-01-03 17:29:14.773869346	INV1003	Tablet	Electronics	2	33618	67236
4	2024-01-04 15:18:59.698492462	INV1004	T-shirt	Fashion	2	4664	9328

```
In [19]: # 1. Overall Sales Summary
        print("Total Revenue:", df["Total_Sales"].sum())
        print("Total Invoices:", df["Invoice_ID"].nunique())
        print("Total Customers:", df["Customer_ID"].nunique())

        # 2. Best Selling Products
        best_products = df.groupby("Product_Name")["Total_Sales"].sum().sort_values(ascending=False)
        print("\nTop 5 Products by Sales:\n", best_products)

        # 3. Category-wise Sales
        category_sales = df.groupby("Category")["Total_Sales"].sum().sort_values(ascending=False)
        print("\nSales by Category:\n", category_sales)

        # 4. Payment Method Analysis
        payment_counts = df["Payment_Method"].value_counts()
```

```

print("\nPayment Method Counts:\n", payment_counts)

# 5. Monthly Sales
monthly_sales = df.groupby("Month")["Total_Sales"].sum().sort_values(ascending=False)
print("\nMonthly Sales:\n", monthly_sales)

# 6. Top Customers
top_customers = df.groupby("Customer_ID")["Total_Sales"].sum().sort_values(ascending=False)
print("\nTop 5 Customers:\n", top_customers)

```

Total Revenue: 7244437

Total Invoices: 200

Total Customers: 49

Top 5 Products by Sales:

Product_Name	Total_Sales
Smartphone	2138093
Laptop	1572252
Tablet	1552704
Camera	1005768
Smartwatch	251131

Name: Total_Sales, dtype: int64

Sales by Category:

Category	Total_Sales
Electronics	6268817
Accessories	490573
Fashion	485047

Name: Total_Sales, dtype: int64

Payment Method Counts:

Payment_Method	count
UPI	57
Debit Card	52
Cash	46
Credit Card	45

Name: count, dtype: int64

Monthly Sales:

Month	Total_Sales
January	1443709
May	1430654
April	1181917
June	1150446
February	1072267
March	965444

Name: Total_Sales, dtype: int64

Top 5 Customers:

Customer_ID	Total_Sales
CUST39	367185
CUST11	364732
CUST18	363814
CUST43	361005
CUST20	345510

Name: Total_Sales, dtype: int64

visualization

Best Selling Products (Bar Chart)


```
In [20]: # Top 5 Products by Sales
best_products = df.groupby("Product_Name")["Total_Sales"].sum().sort_values(ascending=False)

plt.figure(figsize=(8,5))
sns.barplot(x=best_products.values, y=best_products.index, palette="viridis")
plt.title("Top 5 Best Selling Products")
plt.xlabel("Total Sales")
plt.ylabel("Product Name")
plt.show()
```

C:\Users\gonda\AppData\Local\Temp\ipykernel_9912\708407496.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=best_products.values, y=best_products.index, palette="viridis")
```



Category-wise Sales (Bar Chart)


```
In [21]: category_sales = df.groupby("Category")["Total_Sales"].sum().sort_values(ascending=False)

plt.figure(figsize=(6,4))
sns.barplot(x=category_sales.values, y=category_sales.index, palette="magma")
plt.title("Sales by Category")
plt.xlabel("Total Sales")
plt.ylabel("Category")
plt.show()
```

C:\Users\gonda\AppData\Local\Temp\ipykernel_9912\2733312931.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=category_sales.values, y=category_sales.index, palette="magma")
```




Payment Method Distribution (Pie Chart)

```
In [22]: payment_counts = df["Payment_Method"].value_counts()

plt.figure(figsize=(6,6))
plt.pie(payment_counts, labels=payment_counts.index, autopct="%1.1f%%", startangle=90)
plt.title("Payment Method Distribution")
plt.show()
```

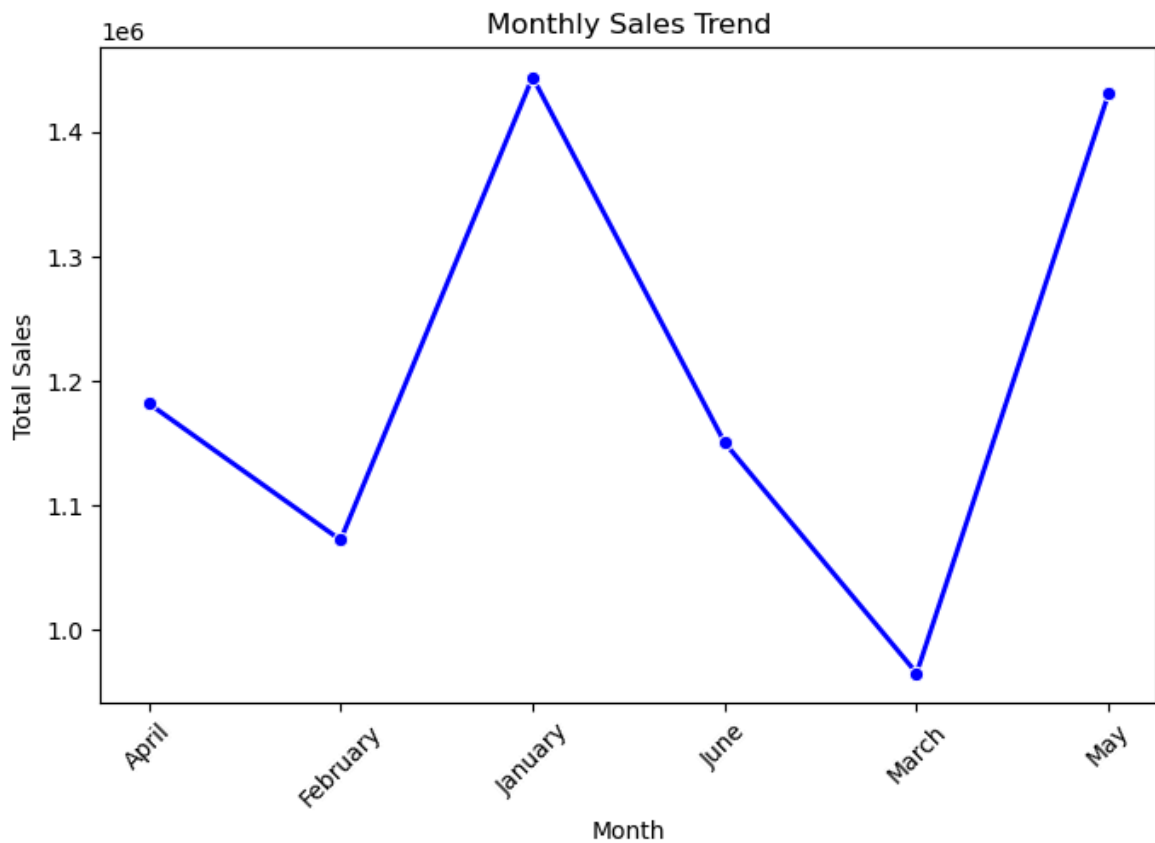
Payment Method Distribution



Monthly Sales Trend (Line Chart)

```
In [23]: monthly_sales = df.groupby("Month")["Total_Sales"].sum()

plt.figure(figsize=(8,5))
sns.lineplot(x=monthly_sales.index, y=monthly_sales.values, marker="o", linewidth=2)
plt.title("Monthly Sales Trend")
plt.xlabel("Month")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.show()
```



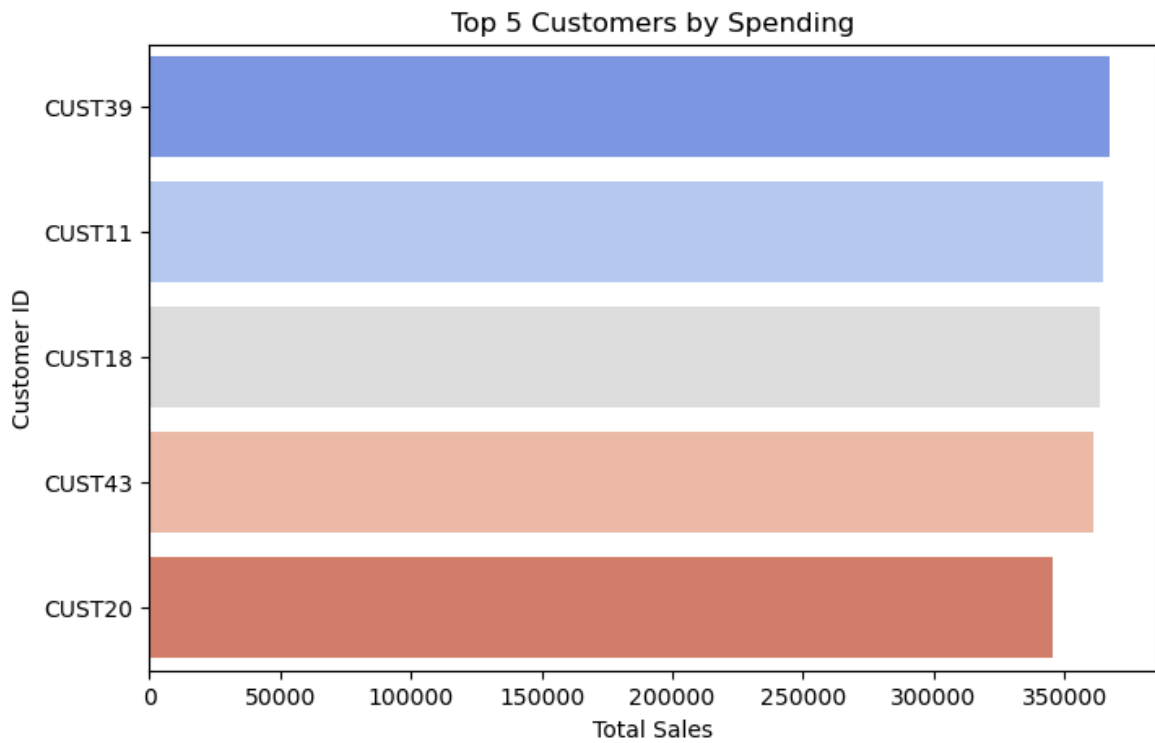
```
In [24]: top_customers = df.groupby("Customer_ID")["Total_Sales"].sum().sort_values(ascen

plt.figure(figsize=(8,5))
sns.barplot(x=top_customers.values, y=top_customers.index, palette="coolwarm")
plt.title("Top 5 Customers by Spending")
plt.xlabel("Total Sales")
plt.ylabel("Customer ID")
plt.show()
```

C:\Users\gonda\AppData\Local\Temp\ipykernel_9912\1553665695.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_customers.values, y=top_customers.index, palette="coolwarm")
```

```
In [27]: # Month wise sales
df['Date'] = pd.to_datetime(df['Date'])
monthly_sales = df.groupby(df['Date'].dt.month)["Total_Sales"].sum().reset_index
monthly_sales.columns = ["Month", "Revenue"]

print(monthly_sales)
```

	Month	Revenue
0	1	1443709
1	2	1072267
2	3	965444
3	4	1181917
4	5	1430654
5	6	1150446

```
In [28]: # Features (X) and Target (y)
X = monthly_sales[["Month"]] # month number as feature
y = monthly_sales["Revenue"]

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluation
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R² Score:", r2_score(y_test, y_pred))
```

MAE: 316995.65

R² Score: -3.3928870786432004

```
In [29]: # Predict for month = 7 (July)
future_month = [[7]]
future_sales = model.predict(future_month)
print("Predicted Sales for July:", future_sales[0])
```

Predicted Sales for July: 1383051.0

C:\Users\gonda\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(

```
In [30]: df.to_csv("cleaned_sales_data.csv", index=False)
```

```
In [ ]:
```